

# Understanding Activation Functions in Deep Learning

A Practical Tutorial on ReLU, Leaky ReLU,  
and ELU with Fashion-MNIST

<b>Author</b>	Md Shahedur Rahman
<b>StudentID</b>	23036883
<b>Course Module</b>	Machine Learning and Neural Networks
<b>Assignment</b>	Individual assignment: Machine learning tutorial

# **Tutorial Outline**

- **Introduction**
- **Dataset Preparation**
- **Building Models with Different Activation Functions**
- **Training and Comparing Models**

# **Tutorial Outline**

- **Advanced Visualization: Decision Boundaries**
- **Observations and Insights**
- **Conclusion**
- **References**

# Introduction

## What Are Activation Functions?

- Activation functions introduce non-linearity into neural networks, enabling them to learn complex patterns (Goodfellow et al., 2016).
- They determine how input signals are transformed into output signals at each layer of the network.

# Introduction

## Why Are Activation Functions Crucial in Deep Learning?

- ❖ Enable networks to approximate non-linear relationships between input and output.
- ❖ Allow stacking of multiple layers to form deep architectures.
- ❖ Examples of activation functions:
  - **ReLU**: Simplicity and efficiency.
  - **Leaky ReLU**: Gradient flow improvements.
  - **ELU**: Smooth convergence for deeper models.

# Understanding Activation Functions

## ReLU: The Most Widely Used Activation Function

- **Definition:**

- Rectified Linear Unit (ReLU) is defined as:

$$f(x) = \max(0, x)$$

- Introduced in the context of deep learning to improve gradient flow and computational efficiency (He et al., 2015).

- **Advantages:**

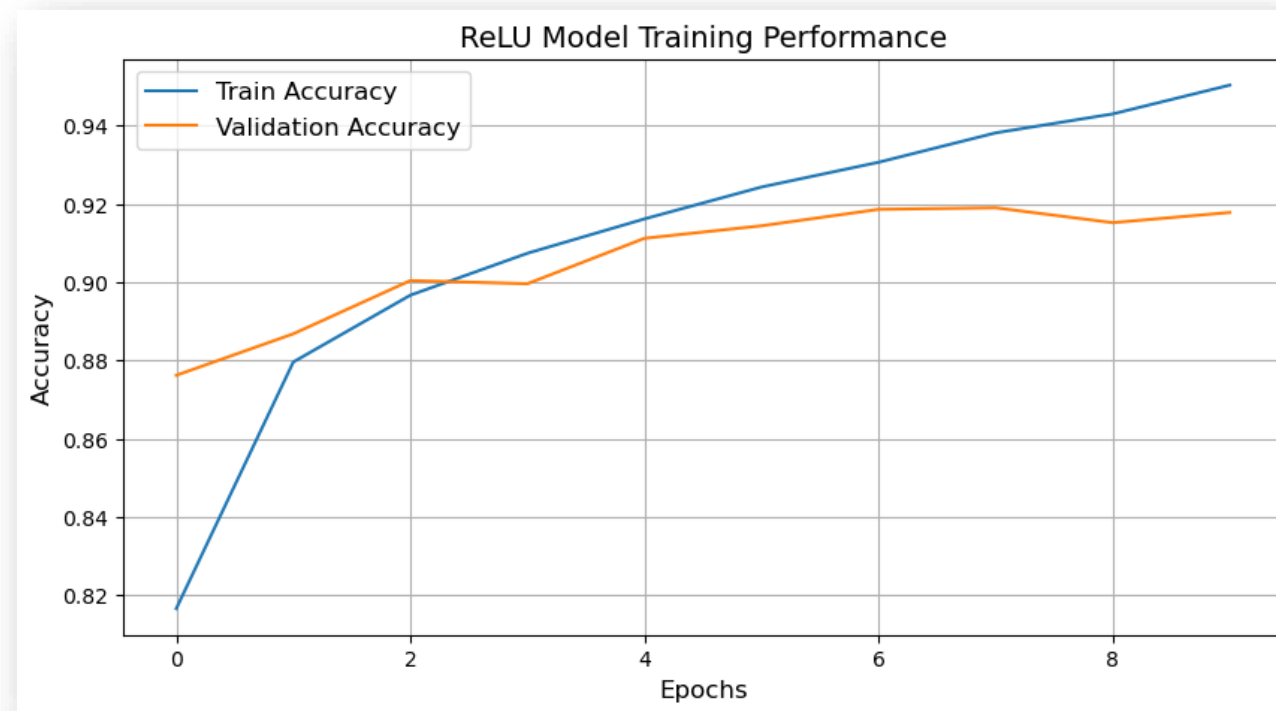
- Computational efficiency.
  - Introduces sparsity into activations.

- **Limitations:**

- Dying ReLU problem: Neurons become inactive for negative inputs.

# Understanding Activation Functions

## ReLU: The Most Widely Used Activation Function





# Understanding Activation Functions

## Leaky ReLU: Addressing the "Dying ReLU" Problem

- **Definition:**

- Leaky ReLU allows a small gradient for negative inputs:  $f(x)=x$  if  $x>0$ , otherwise  $f(x)=\alpha x$ , where  $\alpha$  is a small positive constant (Maas et al., 2013).

- **Advantages:**


- Resolves the dying ReLU issue by allowing gradient flow for negative inputs.

- **Limitations:**

- The small slope in negative regions can still result in slow training in some cases.

# Understanding Activation Functions

## Leaky ReLU: Addressing the "Dying ReLU" Problem



# Understanding Activation Functions

## ELU: Smoothing Convergence for Deep Networks

- **Definition:**

- ELU (Exponential Linear Unit) outputs:  $f(x)=x$  if  $x>0$ , otherwise  $f(x)=\alpha(\exp(x)-1)$  (Clevert et al., 2015).

- **Advantages:**


- Smooth transition between negative and positive inputs improves optimization.

- **Limitations:**

- Slightly more computationally expensive due to the exponential operation.

# Understanding Activation Functions

## ELU: Smoothing Convergence for Deep Networks



# Dataset Preparation

## Introduction to the Fashion-MNIST Dataset

- **Fashion-MNIST:**

- A dataset of 70,000 grayscale images of fashion items, split into 10 classes (e.g., shirts, shoes) (Xiao et al., 2017).
- Challenges: Diverse and ambiguous patterns, making it ideal for testing neural networks.

# Dataset Preparation

## Data Preprocessing Steps

- Normalize pixel values to a range of  $[0, 1]$ .
- Add a channel dimension for compatibility with Conv2D layers.
- Split into training, validation, and test sets.

# Dataset Preparation

## Data Preprocessing Steps

```
# Load Fashion-MNIST dataset
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

# Normalize the data
x_train = x_train / 255.0 # Scale images to [0, 1]
x_test = x_test / 255.0

# Add a channel dimension (required for Conv2D)
x_train = x_train[..., tf.newaxis]
x_test = x_test[..., tf.newaxis]

# Convert labels to categorical (one-hot encoding)
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)

# Split validation data from training
x_val, y_val = x_train[-5000:], y_train[-5000:]
x_train, y_train = x_train[:-5000], y_train[:-5000]

print(f"Training set: {x_train.shape}, Validation set: {x_val.shape}, Test set: {x_test.shape}")
```

# Building Models with Different Activation Functions

## Model Architecture for ReLU

- A simple convolutional neural network with:
  - 2 Conv2D layers using ReLU activation.
  - 2 MaxPooling2D layers for downsampling.
  - 1 Dense layer with 128 units and ReLU.
  - A final Dense layer with softmax for classification.



# Building Models with Different Activation Functions

## Model Architecture for ReLU

```
def create_relu_model():
    model = Sequential([
        Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
        MaxPooling2D((2, 2)),
        Conv2D(64, (3, 3), activation='relu'),
        MaxPooling2D((2, 2)),
        Flatten(),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax') # 10 classes for Fashion-MNIST
    ])
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

relu_model = create_relu_model()
relu_model.summary()
```

# Building Models with Different Activation Functions

## Model Architecture for Leaky ReLU

- Similar to the ReLU model but replaces ReLU with Leaky ReLU:
  - Conv2D layers with Leaky ReLU activation ( $\alpha=0.1$ ).
  - Dense layers with Leaky ReLU activation.

# Building Models with Different Activation Functions

## Model Architecture for Leaky ReLU

```
def create_leaky_relu_model():
    model = Sequential([
        Conv2D(32, (3, 3)),
        LeakyReLU(alpha=0.1), # Leaky ReLU activation
        MaxPooling2D((2, 2)),
        Conv2D(64, (3, 3)),
        LeakyReLU(alpha=0.1),
        MaxPooling2D((2, 2)),
        Flatten(),
        Dense(128),
        LeakyReLU(alpha=0.1),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

leaky_relu_model = create_leaky_relu_model()
```

# Building Models with Different Activation Functions

## Model Architecture for ELU

- Similar to the previous models but replaces activation functions with **ELU**:
  - Conv2D layers with ELU activation.
  - Dense layers with ELU activation for smooth gradients.

# Building Models with Different Activation Functions

## Model Architecture for ELU

```
def create_elu_model():
    model = Sequential([
        Conv2D(32, (3, 3)),
        ELU(alpha=1.0), # ELU activation
        MaxPooling2D((2, 2)),
        Conv2D(64, (3, 3)),
        ELU(alpha=1.0),
        MaxPooling2D((2, 2)),
        Flatten(),
        Dense(128),
        ELU(alpha=1.0),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

elu_model = create_elu_model()
```

# Training Models with Different Activation Functions

- **Training Parameters:**
  - Optimizer: Adam
  - Loss Function: Categorical Crossentropy
  - Batch Size: 64
  - Epochs: 10
- **Trained models:**
  - ReLU Model
  - Leaky ReLU Model
  - ELU Model

# Training Models with Different Activation Functions

```
history_relu = relu_model.fit(  
    x_train, y_train,  
    validation_data=(x_val, y_val),  
    epochs=10,  
    batch_size=64  
)  
  
# Evaluate on test set  
relu_test_loss, relu_test_accuracy = relu_model.evaluate(x_test, y_test)  
print(f"ReLU Model - Test Loss: {relu_test_loss}, Test Accuracy: {relu_test_accuracy}")
```

# Training Models with Different Activation Functions

```
history_leaky_relu = leaky_relu_model.fit(  
    x_train, y_train,  
    validation_data=(x_val, y_val),  
    epochs=10,  
    batch_size=64  
)  
  
history_elu = elu_model.fit(  
    x_train, y_train,  
    validation_data=(x_val, y_val),  
    epochs=10,  
    batch_size=64  
)
```




# Training Results

## Observations:

- All three models converged within 10 epochs.
- Slight differences in validation loss and accuracy, depending on the activation function.


# Training Results

The training and validation accuracy comparison plot



# Training Results

The training and validation loss comparison plot



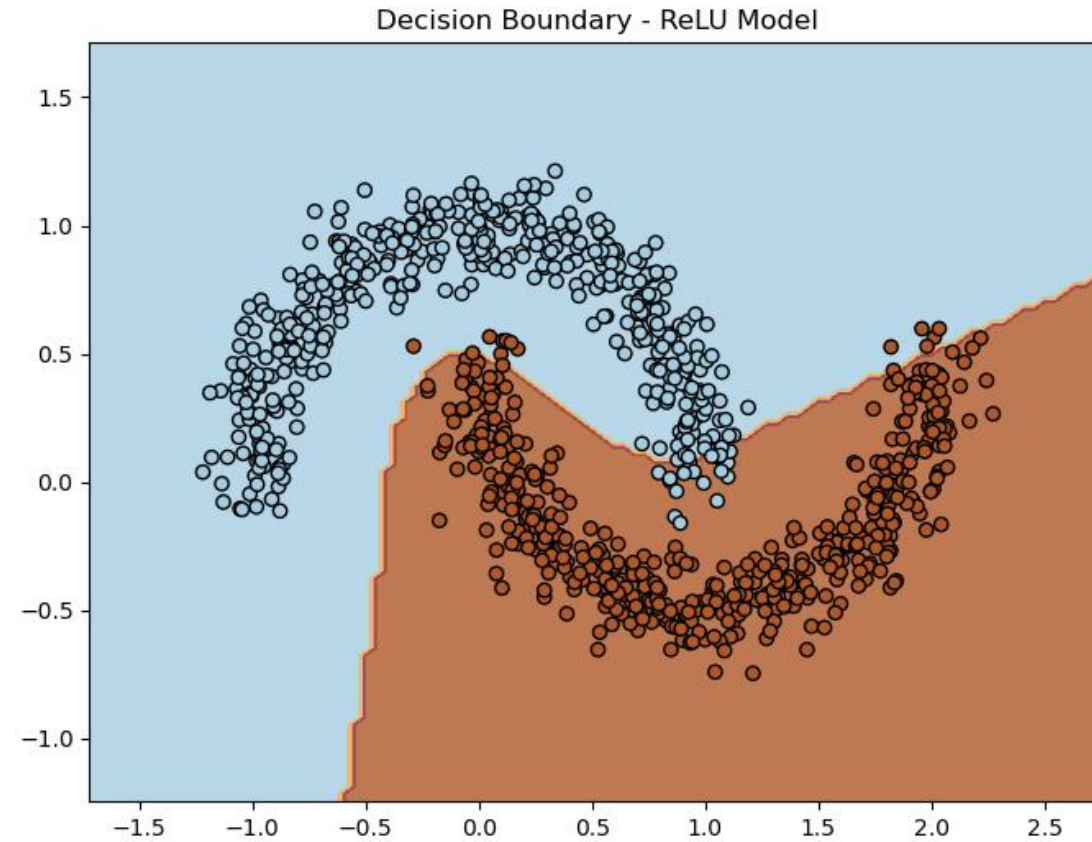
# Decision Boundaries

## Visualizing Decision Boundaries

- Decision boundaries illustrate how models classify data in input space.
- Differences in boundary smoothness reflect the activation functions used.

# Decision Boundaries

## Visualizing Decision Boundaries



# Observations and Insights

- **Performance Comparison:**

- ReLU:

- Fast convergence but prone to dying ReLU for deeper layers.

- Leaky ReLU:

- Improved gradient flow but slower convergence.

- ELU:

- Smooth convergence but higher computational cost.

# Observations and Insights

- **Recommendations:**

- Use ReLU for general-purpose tasks.
- Consider Leaky ReLU for avoiding dying ReLU issues.
- Use ELU for deeper networks requiring smoother optimization.

# Performance Comparison


## Validation Accuracy Comparison

- **ReLU:**
  - Validation accuracy stabilizes around **91%** after 10 epochs.
  - Slight fluctuations in later epochs indicate possible overfitting.
- **Leaky ReLU:**
  - Validation accuracy reaches ~92%, higher than ReLU in the later epochs.
  - Smoothest convergence with reduced overfitting tendencies.
- **ELU:**
  - Validation accuracy starts lower than ReLU and Leaky ReLU but catches up to reach around 91%.
  - Demonstrates smooth convergence but slightly lags in early epochs.



# Performance Comparison

## Validation Accuracy Comparison



# Key Insights

- **When to Use Each Activation Function:**

- ReLU: Best for shallow or moderately deep networks.
- Leaky ReLU: Use when avoiding the "dying ReLU" problem.
- ELU: Ideal for deep networks requiring smoother optimization.

- **Considerations:**

- Computational cost increases from ReLU to ELU.
- Choice of activation can significantly impact training dynamics and accuracy.

# Conclusion

- Activation functions are crucial for introducing non-linearity in neural networks.
- Each activation function (ReLU, Leaky ReLU, ELU) has distinct advantages and trade-offs.
- Optimal selection depends on:
  - Network depth.
  - Task complexity.
  - Desired computational efficiency.

# Reference lists

- Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. MIT Press.
- Maas, A.L., Hannun, A.Y., and Ng, A.Y. (2013). *Rectifier Nonlinearities Improve Neural Network Acoustic Models*. Proceedings of ICML.
- Clevert, D.-A., Unterthiner, T., and Hochreiter, S. (2015). *Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs)*. arXiv:1511.07289.
- Xiao, H., Rasul, K., and Vollgraf, R. (2017). *Fashion-MNIST: A Novel Image Dataset for Benchmarking Machine Learning Algorithms*. arXiv:1708.07747.
- He, K., Zhang, X., Ren, S., and Sun, J. (2015). *Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification*. arXiv:1502.01852.